

Knowledge Cutoff Effects on LLM-Generated Literature Reviews in Rapidly Evolving Fields

A Review of Temporal Limitations, Mitigation Strategies, and Open Research Problems

Abstract

Large language models (LLMs) are increasingly used to accelerate literature review, drafting outlines, synthesizing findings, and generating citations across biomedical, computational, and engineering fields whose publication volume now outpaces manual synthesis capacity. However, every LLM carries a knowledge cutoff, the date beyond which its parametric memory contains no training signal, and this constraint interacts poorly with domains that change on a monthly or weekly cadence. This paper reviews the mechanisms, empirical evidence, and mitigation strategies associated with knowledge cutoff effects on LLM-generated literature reviews. We synthesize evidence from temporal misalignment research, knowledge-editing literature, automated survey-generation systems (e.g., AutoSurvey, SurveyX), and large-scale citation-hallucination audits to characterize three interacting failure modes: stale coverage of recent findings, temporal leakage and forgetting failures under cutoff-simulation prompting, and fabricated or corrupted references. We discuss current mitigation approaches, principally retrieval-augmented generation (RAG) and knowledge editing, and evaluate their partial effectiveness. We then identify research gaps, including the absence of standardized temporal-freshness metrics for review generation, insufficient longitudinal benchmarks in rapidly evolving fields, and unresolved trade-offs between recency and citation reliability. We conclude with recommendations for researchers, tool builders, and journal editors seeking to responsibly deploy LLMs in the literature review pipeline.

Keywords

large language models; knowledge cutoff; literature review generation; retrieval-augmented generation; citation hallucination; temporal misalignment; knowledge editing; scientometrics; automated survey generation

1. Introduction

The volume of scientific output has grown to a point at which manual literature synthesis is straining under its own weight. In computer science alone, more than 4,000 arXiv submissions carried the phrase “Large Language Model” in their title or abstract during just the first four months of 2024, a rate of accumulation that Wang et al. (2024) explicitly cite as motivation for automating survey writing. Analogous growth rates characterize biomedicine, materials science, and climate research, all fields whose foundational facts, benchmark leaderboards, and consensus positions can shift within a single publication cycle. Large language models (LLMs) have been proposed as a partial remedy: systems such as AutoSurvey (Wang et al., 2024), SurveyX (2025), and LLM-based review-generation pipelines built on retrieval-augmented generation (Wu et al., 2025) now draft outlines, summarize clusters of papers, and compose full review sections with citations attached.

Yet every LLM is trained on a fixed corpus assembled up to a specific point in time — its knowledge cutoff — after which the model has no direct exposure to new publications, terminology, or empirical results (Wikipedia contributors, “Knowledge cutoff,” n.d.). This

constraint is not merely a matter of the model being slightly out of date; it interacts with three properties specific to literature review writing. First, review completeness depends on comprehensive coverage of recent work, so any gap between a field’s current state and a model’s training horizon translates directly into missing citations and stale framing. Second, review credibility depends on citation accuracy, and LLMs are demonstrably prone to generating fabricated or corrupted bibliographic records, a problem that recent large-scale audits show is accelerating in the literature itself (Algaba et al., 2026, as reported in the audit of 111 million references). Third, rapidly evolving fields amplify both problems because the rate of genuine conceptual change — new architectures, revised benchmarks, retracted claims — can outpace the interval between model releases.

This paper synthesizes recent empirical and methodological literature to characterize how knowledge cutoffs specifically affect LLM-generated literature reviews in fast-moving research areas, distinguishing this problem from the more general study of LLM temporal misalignment (Luu et al., 2022) and from the general hallucination literature. We organize the discussion around three interacting phenomena — coverage staleness, cutoff-boundary failures, and citation hallucination — and review the principal mitigation strategy, retrieval augmentation, alongside knowledge-editing alternatives. We close by identifying concrete research gaps and proposing directions for benchmark and tooling development.

2. Background and Related Work

2.1 *The Knowledge Cutoff Construct*

A knowledge cutoff is commonly defined as the point in time beyond which a model has not been trained on new data, such that information about later events is simply absent from its parameters unless supplied externally, for example via retrieval-augmented generation (Wikipedia contributors, n.d.). In practice, the notion is less clean than the label suggests. Cheng et al. (2024), in “Dated Data: Tracing Knowledge Cutoffs in Large Language Models,” show that the effective cutoff inferred from a model’s actual behavior frequently diverges from the cutoff date reported by its developer, an inconsistency they trace to two causes: temporal misalignment within CommonCrawl-derived pretraining dumps, which retain substantial quantities of older content, and deduplication procedures that inadvertently discard or conflate near-duplicate documents from different time periods. This finding is significant for literature-review applications because it means a model’s reported cutoff cannot be treated as a reliable guarantee of what the model does or does not “know,” complicating any attempt to reason transparently about coverage gaps.

2.2 *Temporal Misalignment in Language Models*

The broader phenomenon of performance degradation when a model trained on one time period is evaluated on data from another is well established in the temporal misalignment literature. Luu et al. (2022), in “Time Waits for No One!” (NAACL 2022), established a suite of eight tasks across social media, scientific papers, news, and reviews and demonstrated that temporal misalignment measurably degrades end-task performance, with the degree of degradation varying by task. Subsequent work has replicated and extended this finding: Cao et al. (2024), in “Set the Clock,” show a roughly linear decline in factual accuracy as the gap between a model’s effective training period and its evaluation period widens, and related probing work finds that this degradation pattern holds across model scales when models are fine-tuned on yearly data splits. Dhingra et al.

(2022), in “Time-Aware Language Models as Temporal Knowledge Bases” (TACL), propose jointly conditioning language models on textual timestamps to improve both memorization of facts from the training window and calibrated uncertainty about facts from unseen future periods, an early acknowledgment that models need explicit temporal signals rather than an undifferentiated blend of past and present text.

2.3 *Dynamic and Time-Sensitive QA Benchmarks*

A parallel line of work has built benchmarks specifically to probe how LLMs handle evolving facts. RealTime QA (Kasai et al., 2022) delivers weekly quizzes drawn from news websites to test models on current events as they occur. FreshQA (Vu et al., 2023), introduced alongside the FreshPrompt search-augmentation method, provides a human-curated set of roughly 600 questions spanning “never-changing,” “slow-changing,” and “fast-changing” categories, and explicitly demonstrates that closed-book LLMs struggle on fast-changing questions while search-augmented prompting substantially improves accuracy. StreamingQA (Liska et al., 2022) organizes a large question set chronologically against a corpus of timestamped news articles to test adaptation to new knowledge over time. These benchmarks were built for general factual QA rather than for literature review synthesis specifically, but they establish the empirical baseline finding this paper builds on: that LLM factual reliability is a decreasing function of the temporal distance between a fact’s emergence and the model’s effective knowledge horizon.

2.4 *Automated Literature Review and Survey Generation Systems*

Systems that automate literature review composition have matured rapidly since 2024. AutoSurvey (Wang et al., 2024, NeurIPS) introduced a staged pipeline consisting of retrieval-based initial gathering, LLM-generated outlining, parallel subsection drafting by specialized LLM instances, integration and refinement, and iterative citation evaluation, using a retrieval database of roughly 530,000 arXiv computer-science abstracts; the authors frame the entire system as a response to “rapidly evolving fields” in which manual survey writing cannot keep pace. SurveyX (2025) extends this paradigm with an explicit online retrieval module and reports improved citation recall and precision relative to AutoSurvey and naive RAG baselines when benchmarked against human-written surveys. Wu et al. (2025), publishing in National Science Review, describe a domain-general automated review-generation pipeline validated on propane dehydrogenation catalyst literature (343 articles across 35 topics) and report that a statistically validated evaluation framework judges the generated reviews to match or exceed manually written ones on several axes. Complementary systems such as LitLLM combine retrieval-augmented generation with LLM re-ranking against user-supplied abstracts, and ChatCite targets human-like summarization workflows (as summarized in Wu et al., 2025 and related work). Earlier approaches to review automation, by contrast, relied on citation-network structure or multi-document extractive summarization and are noted to “struggle with rapidly evolving fields” and to “underestimate recent publications due to citation lag” (Wu et al., 2025).

2.5 *Citation and Reference Hallucination*

A substantial and rapidly growing literature documents that LLMs fabricate or corrupt bibliographic references even when the surrounding prose is fluent and plausible. Early work by Alkaissi and McFarlane (2023) reported that ChatGPT could produce entirely fictitious but stylistically convincing citations. Tang, Duan, and Cai (2025), in an EMNLP 2025 paper

evaluating reference generation, abstract writing, and review composition, introduce multidimensional metrics for hallucination rate, semantic coverage, and factual consistency, and find that even the most advanced contemporary models still generate hallucinated references, with performance varying substantially by discipline. Agrawal et al. (2024) and a 2026 benchmark of thirteen LLMs (cited within Xu et al., 2026, as reported in the BibTeX Citation Hallucinations survey) find hallucination rates for generated citations ranging from roughly 14% to 95% depending on the vendor and model, with no consistent relationship to batch size. At the scale of the published record itself, a large audit spanning arXiv, bioRxiv, SSRN, and PubMed Central (111 million references across 2.5 million papers) finds a sharp post-2023 rise in non-existent references, with a conservative estimate of 146,932 hallucinated citations appearing in 2025 alone, concentrated in fields with rapid AI-tool uptake and in papers exhibiting linguistic markers of AI-assisted writing. Multiple detection tools — CheckIfExist (Abbonato, 2026), CiteCheck, and HalluCiteChecker — have since been proposed to verify references against CrossRef, Semantic Scholar, and OpenAlex, underscoring that this is now regarded as an operational problem for scholarly publishing rather than a laboratory curiosity. Some commentators go further: a 2026 analysis in a research-ethics journal argues that hallucinated citations can, under some circumstances, constitute research misconduct when citations function evidentially as data within a paper.

3. Main Thematic and Technical Discussion: How Knowledge Cutoffs Distort Review Generation

3.1 Coverage Staleness

The most direct effect of a knowledge cutoff on a literature review is coverage staleness: the systematic under-representation of work published after the model’s effective training horizon. Because an LLM’s parametric knowledge of a topic is proportional to how much relevant text it saw during pretraining, and because that exposure necessarily ends at the cutoff, a closed-book model asked to “summarize recent advances” in a fast-moving subfield will default to describing the state of the field as it existed near the cutoff, often without signaling that this is what it is doing. Wu et al. (2025) explicitly identify citation lag and underestimation of recent publications as a defect of earlier, non-LLM review-generation approaches that relied on citation-network structure, and design their pipeline’s retrieval stage specifically to counteract it. AutoSurvey’s framing likewise treats “rapidly evolving fields” as the central motivating challenge rather than an edge case (Wang et al., 2024).

Coverage staleness compounds with the reported-versus-effective cutoff discrepancy documented by Cheng et al. (2024): if a model’s genuine temporal horizon is earlier than its labeled cutoff, then even a review author who deliberately restricts claims to “pre-cutoff” material may unknowingly omit a further band of legitimate pre-cutoff literature that the model simply never learned well, for instance because it was underrepresented in deduplicated training data.

3.2 Cutoff-Boundary and Forgetting Failures

A second, more subtle failure mode arises when models are explicitly asked to reason as though situated at an earlier point in time — for example, when a review author instructs a model to “describe the state of the art as of 2021” in order to reconstruct historical framing, or when a benchmark author wants to test genuine reasoning rather than memorized post-hoc knowledge.

Gao et al. (2025), in “Can Prompts Rewind Time for LLMs?” construct evaluation sets probing whether prompting can induce an LLM to forget direct factual knowledge, semantic shifts, and causally related knowledge from after a stated cutoff. They find that prompted cutoffs are reasonably effective when the forbidden information is queried directly, but fail when the same information is needed only indirectly, through a causal chain the model does not recognize as connected to the restricted fact. A related 2026 study, “Teaching Large Language Models When Not to Know,” frames this as “ex-ante reasoning” and reports that even leading proprietary models exhibit substantial temporal “leakage” — incorporating knowledge that postdates a stated cutoff — with the best-performing model in their evaluation still leaking over one-fifth of restricted post-cutoff QA content; they further show that standard supervised fine-tuning only partially closes this gap because cutoff-correctness is a relation between an answer and a stated boundary rather than an intrinsic property learnable from static examples, motivating their proposed critique-based fine-tuning framework. A companion paper, “Can LLMs Be Constrained to the Past?,” similarly finds that recall-based prompting reduces but does not eliminate cutoff violations, particularly for facts that require several inferential hops.

These findings matter for literature review generation whenever a review must situate a finding in its correct historical sequence — for instance, correctly describing which architectural idea preceded which benchmark result — because a model prone to temporal leakage may silently anachronize the narrative, attributing knowledge or terminology to an earlier period than it actually existed, or conversely projecting later conceptual frameworks backward onto earlier work in a way that misrepresents the historical record.

3.3 Citation Fabrication Under Recency Pressure

The third and most operationally damaging failure mode is citation hallucination, and there is direct evidence that it interacts with the recency demands of rapidly evolving fields specifically. The large-scale citation audit finds that hallucinated references are disproportionately concentrated in fields experiencing rapid AI-tool uptake, precisely the domains in which authors are most tempted to lean on an LLM to keep pace with the literature, and are more common among small or early-career author teams who may have fewer resources for manual verification. A field-level neuron-tracing study additionally shows that hallucination rates vary systematically by citation style and by the number of references requested in a single generation, evidence that fabrication is not a rare, random event but a structured consequence of how models are prompted for bibliographic content. Because a genuinely novel, recently published paper is by definition absent from a model’s training distribution if it postdates the cutoff, an LLM asked to cite “the most recent work on X” faces a forced choice: it can decline (rare, since instruction-tuned models are optimized toward compliance), it can retrieve the paper via an external tool if one is available, or it can generate a plausible-sounding but non-existent citation that pattern-matches the expected surface form of a recent paper in that subfield. The third outcome is the one repeatedly documented in the literature reviewed in Section 2.5, and it is exactly the outcome most likely under recency pressure in a fast-moving field, since the gap between “what truly is recent” and “what the model was trained on” is largest precisely there.

4. Current Approaches and Developments

4.1 Retrieval-Augmented Generation

The dominant mitigation strategy in both the general LLM literature and the literature-review-generation literature specifically is retrieval-augmented generation (RAG), introduced by Lewis et al. (2020) as a framework combining a pretrained parametric generator with a non-parametric retrieval index, so that generation is grounded in retrieved passages rather than relying solely on memorized weights. In principle, RAG addresses coverage staleness directly: if the retrieval index is kept current, a model can be shown recently published abstracts or full texts at inference time regardless of when its own parameters were frozen. AutoSurvey (Wang et al., 2024) implements this by retrieving from a maintained 530,000-paper arXiv abstract database before drafting; SurveyX layers an online retrieval component on top of a similar pipeline; and Wu et al. (2025) build an end-to-end pipeline from literature retrieval through final review text specifically to avoid the citation-lag problem of earlier, non-retrieval methods. A recent systematic review of 128 highly cited RAG studies (Brown, Roman, and Devereux, 2025) confirms that grounding generation in up-to-date, non-parametric memory while retaining the semantic generalization stored in model weights is now the consensus architectural response to precisely this class of staleness problem, and a complementary systematic review focused on RAG for systematic literature review (SLR) automation proposes a four-stage framework — literature search, screening, data extraction, and synthesis — explicitly designed around the SLR use case (see the 2024 MDPI Applied Sciences review of RAG-based SLR automation).

RAG's effectiveness for review generation is nonetheless partial rather than complete. SurveyX researchers report their citation-quality evaluation using the Citation Recall and Citation Precision metrics originally proposed alongside AutoSurvey, implying that even RAG-based systems must be explicitly audited for whether generated statements are actually supported by retrieved sources and whether cited sources are relevant, rather than assumed correct by construction. Retrieval quality, chunking strategy, and reranking all remain open engineering problems, and RAG does not, by itself, resolve the cutoff-boundary and forgetting failures described in Section 3.2, since a model can still misattribute retrieved information to the wrong time period or blend retrieved and parametric knowledge inconsistently.

4.2 Knowledge Editing

A second, less mature mitigation strategy is knowledge editing: directly modifying a small number of model parameters (locate-then-edit methods such as ROME and MEMIT) or maintaining an explicit external memory of corrected facts, rather than retraining the full model or relying on retrieval at inference time. This literature has begun to grapple explicitly with the temporal dimension of the problem. Yin et al. (2023), introducing the task of Temporal Knowledge Editing and the AToKe benchmark, show that conventional editing methods are effective at inserting new facts but catastrophically forget the historical facts those new facts replaced, and propose a Multi-Editing with Time Objective framework that edits historical and current knowledge jointly so that both remain queryable with correct associated timestamps. Cohen et al. (2024), evaluating the “ripple effects” of knowledge editing in the TACL journal, show that a single edit can leave logically entailed downstream facts inconsistent with the edited fact, a finding directly relevant to literature reviews, where an update to one claim (e.g., a superseded benchmark leaderboard entry)

has ripple implications for surrounding claims (e.g., which method is currently “state of the art”) that editing does not automatically propagate.

4.3 Detection and Verification Tooling

Given that neither RAG nor knowledge editing fully eliminates citation fabrication, a distinct line of tooling has emerged to detect hallucinated references after generation rather than prevent them beforehand. CheckIfExist (Abbonato, 2026) performs multi-source verification of bibliographic references against CrossRef, Semantic Scholar, and OpenAlex. CiteCheck proposes a hybrid, retrieval-grounded framework for detecting hallucinated references specifically in AI-generated scientific reports, motivated explicitly by the observation that LLMs can alter authors, years, titles, DOIs, or URLs while preserving surface plausibility, or fabricate entire non-existent papers. HalluCiteChecker offers a lightweight toolkit oriented toward reducing reviewer workload by flagging candidate hallucinated citations for human verification. Collectively, these tools represent an emerging “verify-after-generate” layer that treats citation hallucination as a problem requiring dedicated infrastructure rather than one that prompting or model choice alone can solve.

5. Research Challenges and Limitations

5.1 The Reported-Versus-Effective Cutoff Problem

As established by Cheng et al. (2024), a model’s documented knowledge cutoff is an unreliable proxy for what the model actually knows, because of temporal contamination in web-scraped corpora and inconsistencies introduced by deduplication. This undermines a natural mitigation strategy — simply telling users “this review may be missing work published after [cutoff date]” — because the true boundary of missing coverage cannot be stated with confidence even by the model’s own developers.

5.2 Imperfect Cutoff-Simulation and Leakage

The ex-ante reasoning literature (the 2026 “Teaching Large Language Models When Not to Know” study and its companion recall-based-prompting paper) demonstrates that prompted or fine-tuned temporal constraints reduce but do not eliminate leakage, and that leakage is worse precisely for causally indirect information — the kind most relevant to synthesizing how a field’s ideas evolved rather than to answering isolated factual queries. This is a structural limitation rather than a prompting bug: because ex-ante correctness depends on the relationship between an answer and an externally specified boundary, and that boundary is not an intrinsic feature of the training data, no purely parametric fix fully resolves it.

5.3 Persistent Citation Hallucination Despite Mitigation

Even with retrieval augmentation, citation hallucination rates remain nontrivial. Tang, Duan, and Cai (2025) report that “even the most advanced models still generate hallucinated references, despite recent progress,” and note substantial cross-disciplinary variation, implying that mitigation strategies validated in one field (e.g., computer science, where arXiv provides a clean, well-indexed retrieval corpus) may transfer poorly to fields with fragmented, paywalled, or less consistently indexed literatures. The field-level hallucination study further shows that hallucination rates vary by citation style and by how many references are requested at once,

suggesting that hallucination is sensitive to surface-level generation conditions in ways that complicate simple mitigation guidance.

5.4 Evaluation and Benchmark Gaps

Much of the temporal misalignment and dynamic-QA literature (RealTime QA, FreshQA, StreamingQA) targets short-answer factual QA rather than long-form review synthesis, so findings about degradation curves and leakage rates cannot be assumed to transfer directly to the review-generation setting without further validation. Review-generation benchmarks that do exist (e.g., the evaluation frameworks accompanying AutoSurvey and SurveyX) emphasize citation recall/precision and content-quality dimensions such as synthesis and critical analysis, but, to date, none of the reviewed systems reports a benchmark specifically measuring temporal freshness — that is, the proportion of a generated review's citations that fall within some recency window appropriate to the field's rate of change, or the rate at which a review's central claims about “current state of the art” have already been superseded by the time of generation.

5.5 Ripple Effects and Internal Consistency

Cohen et al. (2024) show that knowledge edits can leave logically related facts unedited, producing internal inconsistency. Because a literature review is an interconnected narrative in which claims about relative performance, chronological precedence, and methodological lineage all depend on one another, a partial or inconsistent update to a model's knowledge (whether via editing or selective retrieval) risks producing a review that is locally accurate in the sentences directly touched by new information but globally incoherent once those updates are read against the surrounding, unedited narrative.

6. Research Gaps and Future Directions

6.1 Standardized Temporal-Freshness Metrics

Building on the citation recall/precision metrics established by AutoSurvey and extended by SurveyX, future work should develop and standardize temporal-freshness metrics for literature-review generation specifically — for example, a recency-weighted citation coverage score that compares the age distribution of a generated review's citations against the actual publication-date distribution of the field's literature over a matched window, or an explicit “cutoff-disclosure accuracy” metric evaluating whether a system correctly identifies which of its claims are at risk of being outdated.

6.2 Longitudinal Benchmarks in Fast-Moving Fields

Dynamic QA benchmarks such as RealTime QA, FreshQA, and StreamingQA offer a template that has not yet been fully adapted to review-generation evaluation: a benchmark that re-runs the same review-generation prompt at multiple points in time against a field known to be rapidly evolving (e.g., LLM alignment research, single-cell genomics, or battery materials) and measures how quickly generated reviews go stale relative to human-written reviews on the same topic. Such a benchmark would also allow direct measurement of the reported-versus-effective cutoff gap identified by Cheng et al. (2024) in the specific context of review synthesis rather than isolated factual recall.

6.3 Integrating Ex-Ante Reasoning Research into Review Systems

The cutoff-boundary and forgetting-failure findings from the ex-ante reasoning literature have so far been validated primarily on QA-style tasks. An open direction is to test whether critique-based fine-tuning approaches such as TCFT, or recall-based prompting, meaningfully reduce anachronism and leakage specifically within long-form review narratives, where indirect, causally chained information is far more common than in single-turn QA.

6.4 Joint RAG and Knowledge-Editing Pipelines

Because RAG addresses coverage staleness but not internal parametric inconsistency, and knowledge editing addresses targeted factual updates but suffers from ripple-effect incoherence, a promising but underexplored direction is the combination of retrieval-time grounding with editing-time consistency propagation, so that a review-generation system both retrieves current literature and maintains logical consistency among claims that depend on updated facts (for instance, propagating a change in “current best-performing method on benchmark X” to every downstream sentence that references that ranking).

6.5 Discipline-Specific Verification Infrastructure

Given that citation hallucination and retrieval quality both vary by discipline, and that the large-scale citation audit found hallucination concentrated in fields with high AI-tool uptake, future work should extend verification tools such as CheckIfExist, CiteCheck, and HalluCiteChecker beyond their current computer-science-and-general-web orientation into disciplines with less centralized indexing infrastructure (e.g., area studies, applied social science, or clinical subfields not fully covered by PubMed), where the retrieval corpora that RAG-based review generators depend on are inherently less complete.

6.6 Editorial and Community Norms

Finally, given that hallucinated citations have already been identified in accepted conference papers and, in some documented cases, in retracted journal articles, and given the 2026 argument that hallucinated citations functioning as evidentiary data may constitute research misconduct, an important non-technical research direction is the development of disclosure norms and verification requirements for AI-assisted literature reviews at the level of journals, conferences, and funding bodies, paired with the technical detection infrastructure discussed above.

7. Conclusion

Knowledge cutoffs are not a peripheral engineering detail of large language models; they interact directly and consequentially with the task of literature review generation in fields whose factual ground truth changes faster than model release cycles. This review has traced three interacting failure modes documented across a converging body of recent work: coverage staleness, in which recent publications are systematically underrepresented; cutoff-boundary and forgetting failures, in which models leak post-cutoff knowledge or misplace facts in time even under explicit temporal constraints; and citation hallucination, which is empirically worse under the recency pressure characteristic of fast-moving fields and is now measurably contributing to the corruption of the published scientific record. Retrieval-augmented generation remains the most effective and widely deployed mitigation, and dedicated verification tooling has begun to fill the gap that RAG and

prompting alone cannot close, but neither fully resolves the underlying tension between a model's fixed training horizon and a literature that does not stand still. Progress will likely require the joint development of temporal-freshness metrics, longitudinal review-generation benchmarks, consistency-aware editing pipelines, and editorial norms that treat AI-assisted citation verification as a standard, rather than optional, step in scholarly writing.

References

- [1] Abbonato, D. (2026). CheckIfExist: Detecting citation hallucinations in the era of AI-generated content. arXiv. <https://arxiv.org/pdf/2602.15871>
- [2] Alkaissi, H., & McFarlane, S. I. (2023). Artificial hallucinations in ChatGPT: Implications in scientific writing. *Cureus*, 15(2), e35179.
- [3] Anonymous authors. (2026). Automating systematic literature reviews with retrieval-augmented generation: A comprehensive overview. *Applied Sciences*, 14(19), 9103. <https://www.mdpi.com/2076-3417/14/19/9103>
- [4] Anonymous authors. (2026). BibTeX citation hallucinations in scientific publishing agents: Evaluation and mitigation. arXiv:2604.03159. <https://arxiv.org/pdf/2604.03159>
- [5] Anonymous authors. (2026). Can LLMs be constrained to the past? Improving knowledge cutoff through recall-based prompting. arXiv:2606.05804. <https://arxiv.org/html/2606.05804>
- [6] Anonymous authors. (2026). CiteCheck: Retrieval-grounded detection of LLM citation hallucinations in scientific text. arXiv:2605.27700. <https://arxiv.org/pdf/2605.27700>
- [7] Anonymous authors. (2026). Hallucinated citations produced by generative artificial intelligence may constitute research misconduct when citations function as data in scholarly papers. *Accountability in Research*. <https://doi.org/10.1080/08989621.2026.2645390>
- [8] Anonymous authors. (2026). HalluCiteChecker: A lightweight toolkit for hallucinated citation detection and verification in the era of AI scientists. arXiv:2604.26835. <https://arxiv.org/pdf/2604.26835>
- [9] Anonymous authors. (2026). LLM hallucinations in the wild: Large-scale evidence from non-existent citations. arXiv:2605.07723. <https://arxiv.org/abs/2605.07723>
- [10] Anonymous authors. (2026). Teaching large language models when not to know: Learning temporal critique for ex-ante reasoning. arXiv:2605.14636. <https://arxiv.org/pdf/2605.14636>
- [11] Anonymous authors. (2026). Where fake citations are made: Tracing field-level hallucination to specific neurons in LLMs. arXiv:2604.18880. <https://arxiv.org/pdf/2604.18880>
- [12] Attribution in scientific literature: New benchmark and methods (REASONS dataset). (2024). arXiv:2405.02228. <https://arxiv.org/pdf/2405.02228>
- [13] Brown, A., Roman, M., & Devereux, B. (2025). A systematic literature review of retrieval-augmented generation: Techniques, metrics, and challenges. arXiv:2508.06401. <https://arxiv.org/abs/2508.06401>
- [14] Cao, et al. (2024). Set the clock: Temporal alignment of pretrained language models. arXiv:2402.16797. <https://arxiv.org/pdf/2402.16797>
- [15] Cheng, et al. (2024). Dated data: Tracing knowledge cutoffs in large language models. arXiv:2403.12958. <https://arxiv.org/pdf/2403.12958>
- [16] Cohen, R., Biran, E., Yoran, O., Globerson, A., & Geva, M. (2024). Evaluating the ripple effects of knowledge editing in language models. *Transactions of the Association for Computational Linguistics*. https://doi.org/10.1162/tacl_a_00644

- [17] Dhingra, B., Cole, J. R., Eisenschlos, J. M., Gillick, D., Eisenstein, J., & Cohen, W. W. (2022). Time-aware language models as temporal knowledge bases. *Transactions of the Association for Computational Linguistics*, 10, 257–273. https://doi.org/10.1162/tacl_a_00459
- [18] Gao, X., Zhang, R., Du, D., Mahindre, S., Somayajula, S. A., & Xie, P. (2025). Can prompts rewind time for LLMs? Evaluating the effectiveness of prompted knowledge cutoffs. arXiv:2510.02340. <https://arxiv.org/pdf/2510.02340>
- [19] Kasai, J., Sakaguchi, K., Takahashi, Y., Le Bras, R., Asai, A., Yu, X., Radev, D., Smith, N. A., Choi, Y., & Inui, K. (2022). RealTime QA: What's the answer right now? arXiv:2207.13332. <https://arxiv.org/abs/2207.13332>
- [20] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.
- [21] Liska, A., Kocisky, T., Gribovskaya, E., Terzi, T., Sezener, E., Agrawal, D., de Masson d'Autume, C., Scholtes, T., Zaheer, M., Young, S., Gilsenan-McMahon, E., Austin, S., Blunsom, P., & Lazaridou, A. (2022). StreamingQA: A benchmark for adaptation to new knowledge over time in question answering models. *Proceedings of the 39th International Conference on Machine Learning*.
- [22] Luu, K., Khashabi, D., Gururangan, S., Mandyam, K., & Smith, N. A. (2022). Time waits for no one! Analysis and challenges of temporal misalignment. *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 5944–5958. <https://doi.org/10.18653/v1/2022.naacl-main.435>
- [23] SurveyX authors. (2025). SurveyX: Academic survey automation via large language models. arXiv:2502.14776. <https://arxiv.org/html/2502.14776v1>
- [24] Tang, X., Duan, X., & Cai, Z. G. (2025). Large language models for automated literature review: An evaluation of reference generation, abstract writing, and review composition. *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 1602–1617. arXiv:2412.13612.
- [25] Towards mechanistically understanding why memorized knowledge fails to generalize in large language model finetuning. (2026). arXiv:2607.08393. <https://arxiv.org/abs/2607.08393>
- [26] Vu, T., Iyyer, M., Wang, X., Constant, N., Wei, J., Wei, J., Tar, C., Sung, Y.-H., Zhou, D., Le, Q., & Luong, T. (2023). FreshLLMs: Refreshing large language models with search engine augmentation. arXiv:2310.03214. <https://arxiv.org/abs/2310.03214>
- [27] Wang, Y., Guo, Q., Yao, W., Zhang, H., Zhang, X., Wu, Z., Zhang, M., Dai, X., Zhang, M., Wen, Q., Ye, W., Zhang, S., & Zhang, Y. (2024). AutoSurvey: Large language models can automatically write surveys. *Advances in Neural Information Processing Systems*, 37, 115119–115145. arXiv:2406.10252.
- [28] Wikipedia contributors. (n.d.). Knowledge cutoff. Wikipedia. https://en.wikipedia.org/wiki/Knowledge_cutoff
- [29] Wu, S., Ma, X., Luo, D., Li, L., Shi, X., Chang, X., Lin, X., Luo, R., Pei, C., Du, C., Zhao, Z.-J., & Gong, J. (2025). Automated literature research and review-generation method based on large language models. *National Science Review*. <https://doi.org/10.1093/nsr/nwaf169>

- [30] Yin, X., Jiang, J., Yang, L., & Wan, X. (2023). History matters: Temporal knowledge editing in large language model. arXiv:2312.05497. <https://arxiv.org/abs/2312.05497>